

PULP
Parallel Ultra Low Power

PULP: An Open-Source RISC-V ased Multi-Core Platform for In-Sensor Analytics

Workshop on Open Source Design Automation (OSDA) 2019

29.03.2019

Davide Rossi



¹Department of Electrical, Electronic
and Information Engineering

<http://pulp-platform.org>

ETH zürich

²Integrated Systems Laboratory

Parallel Ultra Low Power (PULP)

- Project started in **2013**
- A collaboration between University of Bologna and ETH Zürich
 - Large team. In total we are about 60 people, not all are working on PULP
- Key goal is

**How to get the most BANG
for the ENERGY consumed
in a computing system**

- We were able to start with a clean slate, no need to remain compatible to legacy systems.

How we started with open source processors

- Our research was not developing processors...
- ... but we needed good processors for systems we build for research
- **Initially (2013) our options were**
 - Build our own (support for SW and tools)
 - Use a commercial processor (licensing, collaboration issues)
 - Use what is openly available (OpenRISC,..)
- **We started with OpenRISC**
 - First chips until mid-2016 were all using OpenRISC cores
 - We spent time improving the microarchitecture
- **Moved to RISC-V later**
 - Larger community, more momentum
 - Transition was relatively simple (new decoder)



We have developed several optimized RISC-V cores

RISC-V Cores

RI5CY

32b

Micro
riscy

32b

Zero
riscy

32b

Ariane

64b

Only processing cores are not enough, we need more

RISC-V Cores

RI5CY
32b

Micro
riscy
32b

Zero
riscy
32b

Ariane
64b

Peripherals

JTAG

SPI

UART

I2S

DMA

GPIO

Interconnect

Logarithmic interconnect

APB – Peripheral Bus

AXI4 – Interconnect

Accelerators

HWCE
(convolution)

Neurostream
(ML)

HWCrypt
(crypto)

PULPO
(1st order opt)

All these components are combined into platforms

RISC-V Cores

RI5CY

32b

Micro riscy

32b

Zero riscy

32b

Ariane

64b

Peripherals

JTAG

UART

DMA

SPI

I2S

GPIO


Interconnect

Logarithmic interconnect

APB – Peripheral Bus


AXI4 – Interconnect

Platforms




Single Core

- PULPino
- PULPissimo



Multi-core

- Fulmine
- Mr. Wolf



Multi-cluster

- Hero

IOT

HPC

Accelerators

HWCE
(convolution)

Neurostream
(ML)

HWCrypt
(crypto)

PULPO
(1st order opt)

RISC-V Instruction Set Architecture

- Started by UC-Berkeley in 2010
- RISC-V is an open standard governed by RISC-V foundation
 - Necessary for the continuity
 - Extensions are still being developed
- Defines 32, 64 and 128 bit ISA
 - No implementation, just the ISA
 - Different RISC-V implementations (both open and close source) are available
- The PULP project specializes in **efficient implementations of RISC-V cores and peripherals**


Spec separated into “extensions”

- I** Integer instructions
- E** Reduced number of registers
- M** Multiplication and Division
- A** Atomic instructions
- F** Single-Precision Floating-Point
- D** Double-Precision Floating-Point
- C** Compressed Instructions
- X** Non Standard Extensions

Our RISC-V family explained

32 bit				64 bit
Low Cost Core	Core with DSP enhancements	Floating-point capable Core		Linux capable Core
<ul style="list-style-type: none">■ Zero-riscy■ RV32-ICM <ul style="list-style-type: none">■ Micro-riscy■ RV32-CE	<ul style="list-style-type: none">■ RI5CY<ul style="list-style-type: none">■ RV32-ICMX<ul style="list-style-type: none">■ SIMD■ HW loops■ Bit manipulation■ Fixed point	<ul style="list-style-type: none">■ RI5CY+FPU<ul style="list-style-type: none">■ RV32-ICMFX		<ul style="list-style-type: none">■ Ariane■ RV64-IMAFDCX■ Full privilege specification
ARM Cortex-M0+	ARM Cortex-M4	ARM Cortex-M4F		ARM Cortex-A55

RI5CY – Our workhorse 32-bit core



- 4-stage pipeline, optimized for energy efficiency
- 40 kGE, 30 logic levels, Coremark/MHZ 3.19
- Includes various extensions (Xpulp) to RISC-V for DSP applications

Our extensions to RI5CY (with additions to GCC)

- **Post-incrementing** load/store instructions
- **Hardware Loops** (`lp.start`, `lp.end`, `lp.count`)
- **ALU instructions**
 - Bit manipulation (count, set, clear, leading bit detection)
 - Fused operations: (add/sub-shift)
 - Immediate branch instructions
- **Multiply Accumulate** (32x32 bit and 16x16 bit)
- **SIMD instructions** (2x16 bit or 4x8 bit) with scalar replication option
 - add, min/max, dotproduct, shuffle, pack (copy), vector comparison


For 8-bit values the following can be executed in a single cycle (`pv.dotup.b`)

$$Z = D_1 \times K_1 + D_2 \times K_2 + D_3 \times K_3 + D_4 \times K_4$$

Why we designed other 32-bit cores after RI5CY?


- **RI5CY was built for energy efficiency for DSP applications**
 - Ideally all parts of the core are running all the time doing something useful
 - This does not always mean it is low-power
 - The core is rather large (> 40 kGE without FPU)
- **People asked us about a simple and small core**
 - Not all processor cores are used for DSP applications
 - The DSP extensions are mostly idle for control applications
 - **Zero-Riscy** was designed to as a **simple and efficient** core.
- **Some people wanted the smallest possible RISC-V core**
 - It is possible to further reduce area by using 16 registers instead of 32 (E)
 - Also the multiplier can be removed saving a bit more
 - **Micro-Riscy** is a parametrized variation of Zero-Riscy with **minimal area**

Zero/Micro-riscy, small area core for control applications



- Only 2-stage pipeline, simplified register file
- **Zero-Riscy** (RV32-ICM), 19kGE, 2.44 Coremark/MHz
- **Micro-Riscy** (RV32-EC), 12kGE, 0.91 Coremark/MHz
- Used as SoC level controller in newer PULP systems

Different 32-bit cores with different area requirements



RI5CY


Zero-riscy

Micro-riscy

Finally the step into 64-bit cores



- **For the first 4 years of the PULP project we used only 32bit cores**
 - Most IoT applications work well with 32bit cores.
 - A typical 64bit core is much more than 2x the size of a 32bit core.
- **But times change:**
 - Using a 64bit Linux capable core allows you to share the same address space as main stream processors.
 - We are involved in several projects where we (are planning to) use this capability
 - There is a lot of interest in the security community for working on a contemporary open source 64bit core.
 - Open research questions on how to build systems with multiple cores.

ARIANE: Our Linux Capable 64-bit core




Main properties of Ariane

- Tuned for high frequency, 6 stage pipeline, integrated cache
 - In order issue, out-of-order write-back, in-order-commit
 - Supports privilege spec 1.11, M, S and U modes
 - Hardware Page Table Walker
- Implemented in GF22nm (Poseidon, Kosmodrom), and UMC65 (Scarabaeus)
 - In 22nm: ~1 GHz worst case conditions (SSG, 125/-40C, 0.72V)
 - 8-way 32kByte Data cache and 4-way 32kByte Instruction Cache
 - Core area: 175 kGE




Ariane booting Linux on a Digilent Genesys 2 board




What About Floating Point Support?

- F (single precision) and D (double precision) extension in RISC-V
- Uses separate floating point register file
 - specialized float loads (also compressed)
 - float moves from/to integer register file
- Fully IEEE compliant
- RI5CY support for F
- Ariane for F and D
- Alternative FP Format support (<32 bit)



Packed-SIMD support for all formats




Unified FP/Integer register file

- Not standard
- up to 15 % better performance
 - Re-use integer load/stores (post incrementing Id/st)
 - Less area overhead
 - Useful if pressure on register file is not very high (true for a lot of applications)

Parametric Floating-Point Unit for Transprecision

- Main FP operation groups
 - MUL/ADD: Add/Subtract, Multiply, FMA
 - CMP/SGNJ: Comparisons, Min/Max etc.
 - CAST: FP-FP casts, Int-FP / FP-Int casts
- Parametrizable
 - Number & Encoding of **Formats**
 - Packed-SIMD **Vectors**
 - **# Pipeline Stages** (per Op and Format)
 - **Implementation** (per Op and Format)
 - **PARALLEL** for best Speed
 - **MERGED** (or Iterative) for best Area
- Special Functions for Transprecision
 - Cast-and-Pack 2 FP Values to Vector
 - Casts amongst FP Vectors + Repacking
 - Expanding FMA (e.g. FP32 += FP16*FP16)



Memory Protection (PMU and MMU)

Physical Memory Protection (PMP)

- Protect the physical memory when the core runs in U or S privilege level
- Up to 16 entries for address filtering
- Configuration held in 4 CSRs
pmpcfg[0-3]
- Whether Store (W), Load (R) and Fetch (X) is allowed
- Address matching modes:
 - Naturally aligned power-of-2 regions (NAPOT) or aligned 4 Byte (NA4)
 - Boundaries >, < (TOR)
- Implemented in **RISCY**

Supervisor Memory Translation and Protection (for Linux-like systems)

- Effectively needs TLBs
- Register to configure base page number (**satp**)
- Translation Mode (32, 39, 48 virtual addressing)
- Address Space Identifier (**ASID**)
- Implemented in **Ariane**



The pulp-platforms put everything together

RISC-V Cores				Peripherals		Interconnect	
RI5CY 32b	Micro riscy 32b	Zero riscy 32b	Ariane 64b	JTAG	SPI	Logarithmic interconnect	
				UART	I2S	APB – Peripheral Bus	
				DMA	GPIO	AXI4 – Interconnect	

Platforms



Single Core

- PULPino
- PULPissimo

Accelerators

HWCE
(convolution)


Neurostream
(ML)

HWCrypt
(crypto)

PULPO
(1st order opt)


PULPino our first single core platform

- **Simple design**
 - Meant as a quick release
- **Separate Data and Instruction memory**
 - Makes it easy in HW
 - Not meant as a Harvard arch.
- **Can be configured to work with all our 32bit cores**
 - RI5CY, Zero/Micro-Riscy
- **Peripherals copied from its larger brothers**
 - Any AXI and APB peripherals could be used




PULPissimo the improved single core platform

- **Shared memory**
 - Unified Data/Instruction Memory
 - Uses the multi-core infrastructure
- **Support for Accelerators**
 - Direct shared memory access
 - Programmed through APB bus
 - Number of TCDM access ports determines max. throughput
- **uDMA for I/O subsystem**
 - Can copy data directly from I/O to memory without involving the core
- **Used as a fabric controller in larger PULP systems**




Kerbin the single core support structure for Ariane

- Still work in progress
 - Current version is very simple
 - Useful for in-house testing
 - A more advanced version will likely be developed soon



OpenPiton and Ariane together, the many-core system

- **OpenPiton**
 - Developed by Princeton
 - Originally OpenSPARC T1
 - Scalable NoC with coherent LLC
 - Tiled Architecture
- **Still work in progress**
 - Bare-metal released in Dec '18
 - Update with support for SMP Linux will be released soon



The main PULP systems we develop are cluster based

RISC-V Cores

RI5CY
32b

Micro
riscy
32b

Zero
riscy
32b

Ariane
64b

Peripherals

JTAG

SPI

UART

I2S

DMA

GPIO

Interconnect

Logarithmic interconnect

APB – Peripheral Bus


AXI4 – Interconnect

Platforms



Single Core

- PULPino
- PULPissimo



Multi-core

- Fulmine
- Mr. Wolf

Accelerators

HWCE
(convolution)

Neurostream
(ML)


HWCrypt
(crypto)

PULPO
(1st order opt)


The main components of a PULP cluster

- **Multiple RISC-V cores**
 - Individual cores can be started/stopped with little overhead
 - DSP extensions in cores
- **Multi-banked scratchpad memory (TCDM)**
 - **Not a cache**, there is no L1 data cache in our systems
- **Logarithmic Interconnect allowing all cores to access all banks**
 - Cores will be stalled during contention, includes arbitration
- **DMA engine to copy data to and from TCDM**
 - Data in TCDM managed by software
 - Multiple channels, allows pipelined operation
- **Hardware accelerators with direct access to TCDM**
 - No data copies necessary between cores and accelerators.


PULP cluster contains multiple RISC-V cores




All cores can access all memory banks in the cluster




Data is copied from a higher level through DMA




There is a (shared) instruction cache that fetches from L2




Hardware Accelerators can be added to the cluster



Event unit to manage resources (fast sleep/wakeup)



An additional microcontroller system (PULPissimo) for I/O



Finally multi-cluster PULP systems for HPC applications

RISC-V Cores

RI5CY

32b

Micro riscy

32b

Zero riscy

32b

Ariane

64b

Peripherals

JTAG

UART

DMA

SPI

I2S

GPIO


Interconnect

Logarithmic interconnect

APB – Peripheral Bus


AXI4 – Interconnect

Platforms




Single Core

- PULPino
- PULPissimo



Multi-core

- Fulmine
- Mr. Wolf



Multi-cluster

- Hero

IOT

HPC

Accelerators

HWCE
(convolution)

Neurostream
(ML)


HWCrypt
(crypto)

PULPO
(1st order opt)

Heterogenous Research Platform

industry-standard, hard-macro
Arm Cortex-A Host processor

scalable, configurable, modifiable FPGA implementation
of PULP (silicon-proven, cluster-based PMCA with RISC-V PEs)




- First released in 2018
- Allows a PULP cluster to be connected to a host system

How to accelerate processing in PULP systems

- **Standard peripheral talking over AXI/APB**
 - Common in practice, nothing new or special
- **Instruction set extensions**
 - Possible in RISC-V, RI5CY has many new DSP instructions
 - Goal is to get ‘**good instructions**’ into standard extensions at some point
- **Shared functional units**
 - Amortizes expensive extensions (FPU/DIV) between multiple units
- **Additional pipeline stages**
 - Work in Patronus for Control flow Integrity
- **Shared memory accelerators**
 - Our bread and butter, PULPopen, NTX
- **Cluster as an accelerator**
 - HERO, BigPULP, etc

What Kind of Acceleration: ISA Extensions

- High Flexibility, relatively small performance boost
- Integrated in the Pipeline of Processors (ID Stage, EX Stage, WB Stage)
- Suffer from Register File Bandwidth Bottleneck (only 2 operands...)
- Require Adapting Compiler and Binutils
- Auxiliary Processing Units (APU), interface available in the RI5CY
- Examples: Dot product (already implemented), bit-reverse, butterfly...




Programming model

```
#define SumDotp(a, b, c) \
__builtin_pulp_sdotsp2(a, b, c)

for (int k = 0; k < (N>>1); k++) {
    VA = VectInA[k];
    VB = VectInB[k];
    S = SumDotp(VA, VB, S);
}
```

What Kind of Acceleration: Shared functional units


- The same as previous one, but one unit can be shared among multiple cores
- Useful to save area for low-utilization instructions (i.e. $< 1/\#cores\%$)
- Examples: Shared FPU (SQRT, DIV...)




What Kind of Acceleration: Shared memory accelerators

Coarse-Grained Shared-Memory Accelerators

- DFGs mapped In Hardware (ILP + DLP) → Highest Efficiency, Low Flexibility
- Sharing data memory with processor for fast communication → low overhead
- Controlled through a memory-mapped interface
- Typically one/two accelerators shared by multiple cores




Fulmine: a HW-Accelerated Secure IoT System-on-Chip



- UMC 65nm technology
 - 6.86 mm²
- 4 cores, 2 accelerators
 - HWCE for 3D conv layers
 - DSP-optimized cores
 - **HWCrypt for AES**
- 64 kB of L1, 192 kB of L2
- First version of uDMA for I/O with no SW intervention
 - QSPI master/slave
 - I2C
 - I2S
 - UART


The Quicklogic eFPGA integration in PULPissimo



- APB port for configuration, programming and control
- Direct 4x TCDM access for eFPGA
 - 128 bits/cycle
 - 4 independent R/W
- Possibility to use uDMA


Blocks not drawn to scale

What Kind of Acceleration: Cluster as an accelerator



Cluster of
RISC-V cores
as an accelerator

We have designed more than 25 ASICs based on PULP




ASICs meant to go on IC Tester




- Mainly characterization
- Not so many peripherals

ASICs meant for applications

- More peripherals (SPI, Camera)
- More on-chip memory





Brief illustrated history of selected ASICs



- All are 28 FDSOI technology, RVT, LVT and RVT flavor
- Uses OpenRISC cores
- Chips designed in collaboration with STM, EPFL, CEA/LETI
- PULPv3 has ABB control

The first system chips, meant for applications




- First multi-core systems that were designed to work on development boards. Each have several peripherals (SPI, I2C, GPIO)
- **Mia Wallace** and **Fulmine** (UMC65) use OpenRISC cores
- **Honey Bunny** (GF28 SLP) uses RISC-V cores
- All chips also have our own FLL designs.

Combining PULP with analog front-end for Biomedical apps




- Designed in collaboration with the Analog group of Prof. Huang
- All chips with SMIC130 (because of analog IPs)
- First three with OpenRISC, VivoSoC3 with RISC-V

The new generation chips from 2018



- System chips in TSMC40 (Mr. Wolf) and UMC65
- **Mr. Wolf:** IoT Processor with 9 RISC-V cores (Zero-riscy + 8x RI5CY)
- **Atomario:** Multi cluster PULP (2x clusters with 4x RI5CY cores each)
- **Scarabaeus:** Ariane based microcontroller

The large system chips from 2018





- All are 22nm Globalfoundries FDX, around 10 sqmm, 50-100 Mtrans
- **Poseidon:** PULPissimo (RI5CY) + Ariane
- **Kosmodrom:** 2x Ariane + NTX (due this week)
- **Arnold:** PULPissimo (RI5CY) + Quicklogic eFPGA

You can buy development boards with PULP technology

VEGA board from open-isa.org

- Micro-controller board with RI5CY and zero-riscy



GAPUINO from Greenwaves

- PULP cluster system with Nine RI5CY cores

We firmly believe in Open Source movement

The screenshot shows the PULP Platform website at <https://www.pulp-platform.org>. The header includes navigation links for PULP Platform, Resources, About, FAQ, Privacy Policy, and Contact. A search bar is also present. The main content features a background image of several PULP SoCs and the text "PULP Platform" and "Open hardware, the way it should be!". A countdown timer indicates "104 days left until the RISC-V conference at ETH Zurich". Below this, a diagram illustrates the PULP architecture components: RISC-V Cores (RI5CY, Micro riscy, Zero riscy, Ariane), Peripherals (JTAG, SPI, UART, I2S, DMA, GPIO), and Interconnect (Logarithmic interconnect, APB – Peripheral Bus, AXI4 – Interconnect). It also shows the evolution from Single Core (PULPino, PULPissimo) to Multi-core (Fulmine, Mr. Wolf) and Multi-cluster (Hero) systems, along with IOT and HPC applications. Social media links for GitHub and Twitter are provided.

First launched in February 2016 (Github)

All our development is on open repositories

Contributions from many groups

We provide PULP with SOLDER Pad License

- Similar to Apache/BSD, adapted specifically for Hardware
- Allows you to:
 - Use
 - Modify
 - Make products and sell them without restrictions.
- Note the difference to **GPL**
 - Systems that include PULP do not have to be open source (Copyright not Copyleft)
 - They can be released commercially
 - LGPL may not work as you think for HW

SOLDER Pad

<http://www.solderpad.org/licenses/>

Silicon and Open Hardware fuel PULP success

- Many companies (we know of) are actively using PULP
 - They value that it is **silicon proven**
 - They like that it uses a **permissive open source license**



Direct research collaborators on PULP

Politecnico di Torino	
University of Cambridge	
USI Lugano	
TU Kaiserslautern	
University of Cagliari	

IBM Research Zurich	
EPF Lausanne	
CSEM Neuchatel	
Princeton University	

Technische Universität Graz	
CEA-Leti Grenoble	
Fraunhofer-Gesellschaft	
Sapienza Università di Roma	

Academic users we are aware of

Università di Genova	
Politecnico di Milano	
Fondazione Bruno Kessler	
Lund University	

Stanford University	
UC Los Angeles	
UC San Diego	
Columbia University	

Universitat Bar-Ilan	
İstanbul Teknik Üniversitesi	
NCTU Hsinchu	
University of Zagreb, FER	

TUT Tampere	
RWTH Aachen	
IST University of Lisboa	
UFRN Rio Grande do Norte	

TU Darmstadt	
Universität Bremen	
Hongik University Seoul	
IIT Kharagpur	

<http://pulp-platform/wosh>

WOSH

the Week of Open Source Hardware

June 11-14 Zürich, SWITZERLAND

- Official RISC-V Workshop (June 11-12)
- RISC-V foundation member meetings (June 13)
- Eurolab4HPC, Open Source Innovation Camp (June 13)
- Licensing and IP rights for Open source HW (June 13)
- FOSSI: Path to high quality IP, Open source EDA tools (June 14)
- Tutorials, demos, hackathons


Week of Open Source Hardware, June 11-14, Zurich

	 RISC-V Paid RISC-V FOUNDATION	 FOSSI Foundation	 Eurolab FREE WOSH SESSIONS	Registration (expect early March)
Tuesday 11 th of June	RISC-V Day1 OPEN Now			Working on Program Now
Wednesday 12 th of June	RISC-V Day2			
Thursday 13 th of June	Member meetings	Licensing Innovation with WOSH	PULP OpenRISC	If you have ideas Contact me
Friday 14 th of June		High Quality IP in DASH Open Source EDA tools	MyriadRF OpenISA.org Greenwaves Hero	

Future of PULP, what we expect

- **Working on a non-profit organization to:**
 - Manage the distribution of PULP
 - Governance
 - Technical Support
 - Continuity
- **ETH Zürich and University of Bologna will be contributors**
 - We can concentrate on our research, energy efficient processor systems
 - Continue using PULP in our work
 - Become a regular contributor
- **Hope to announce during the WOSH in Zurich**
 - There is some work towards this goal, nothing official yet

QUESTIONS?



@pulp_platform
<http://pulp-platform.org>